



Section with Barbara

Week 4


AI Stories




Projects



Learning



AI Stories: AlphaGo



SUPERVISED LEARNING

Make predictions
from data

UNSUPERVISED LEARNING

Find patterns
in data

REINFORCEMENT LEARNING

Generates data from an
environment

Reinforcement Learning

An Introduction
second edition

Richard S. Sutton and Andrew G. Barto




O'REILLY®

Hands-on Machine Learning with Scikit-Learn, Keras & TensorFlow

Concepts, Tools, and Techniques to Build Intelligent Systems


powered by
jupyter



Aurélien Géron

2nd Edition
Updated for
TensorFlow 2

O'REILLY®



powered by
jupyter

Jake VanderPlas

Springer Series in Statistics

Trevor Hastie
Robert Tibshirani
Jerome Friedman

The Elements of Statistical Learning

Data Mining, Inference, and Prediction

Second Edition

Springer

PROCESS

Data Collection, Feature Engineering, Server Infrastructure

PAYOFF

Configuration, Evaluation, Verification, Debugging

PREDICTION

Automation, Deployment, Management, Monitoring

PROCESS

SQL, JSON, Pandas, Numpy, APIs, etc

PAYOFF

Sci-Kit Learn, Keras, Tensorflow, Pytorch, etc

PREDICTION

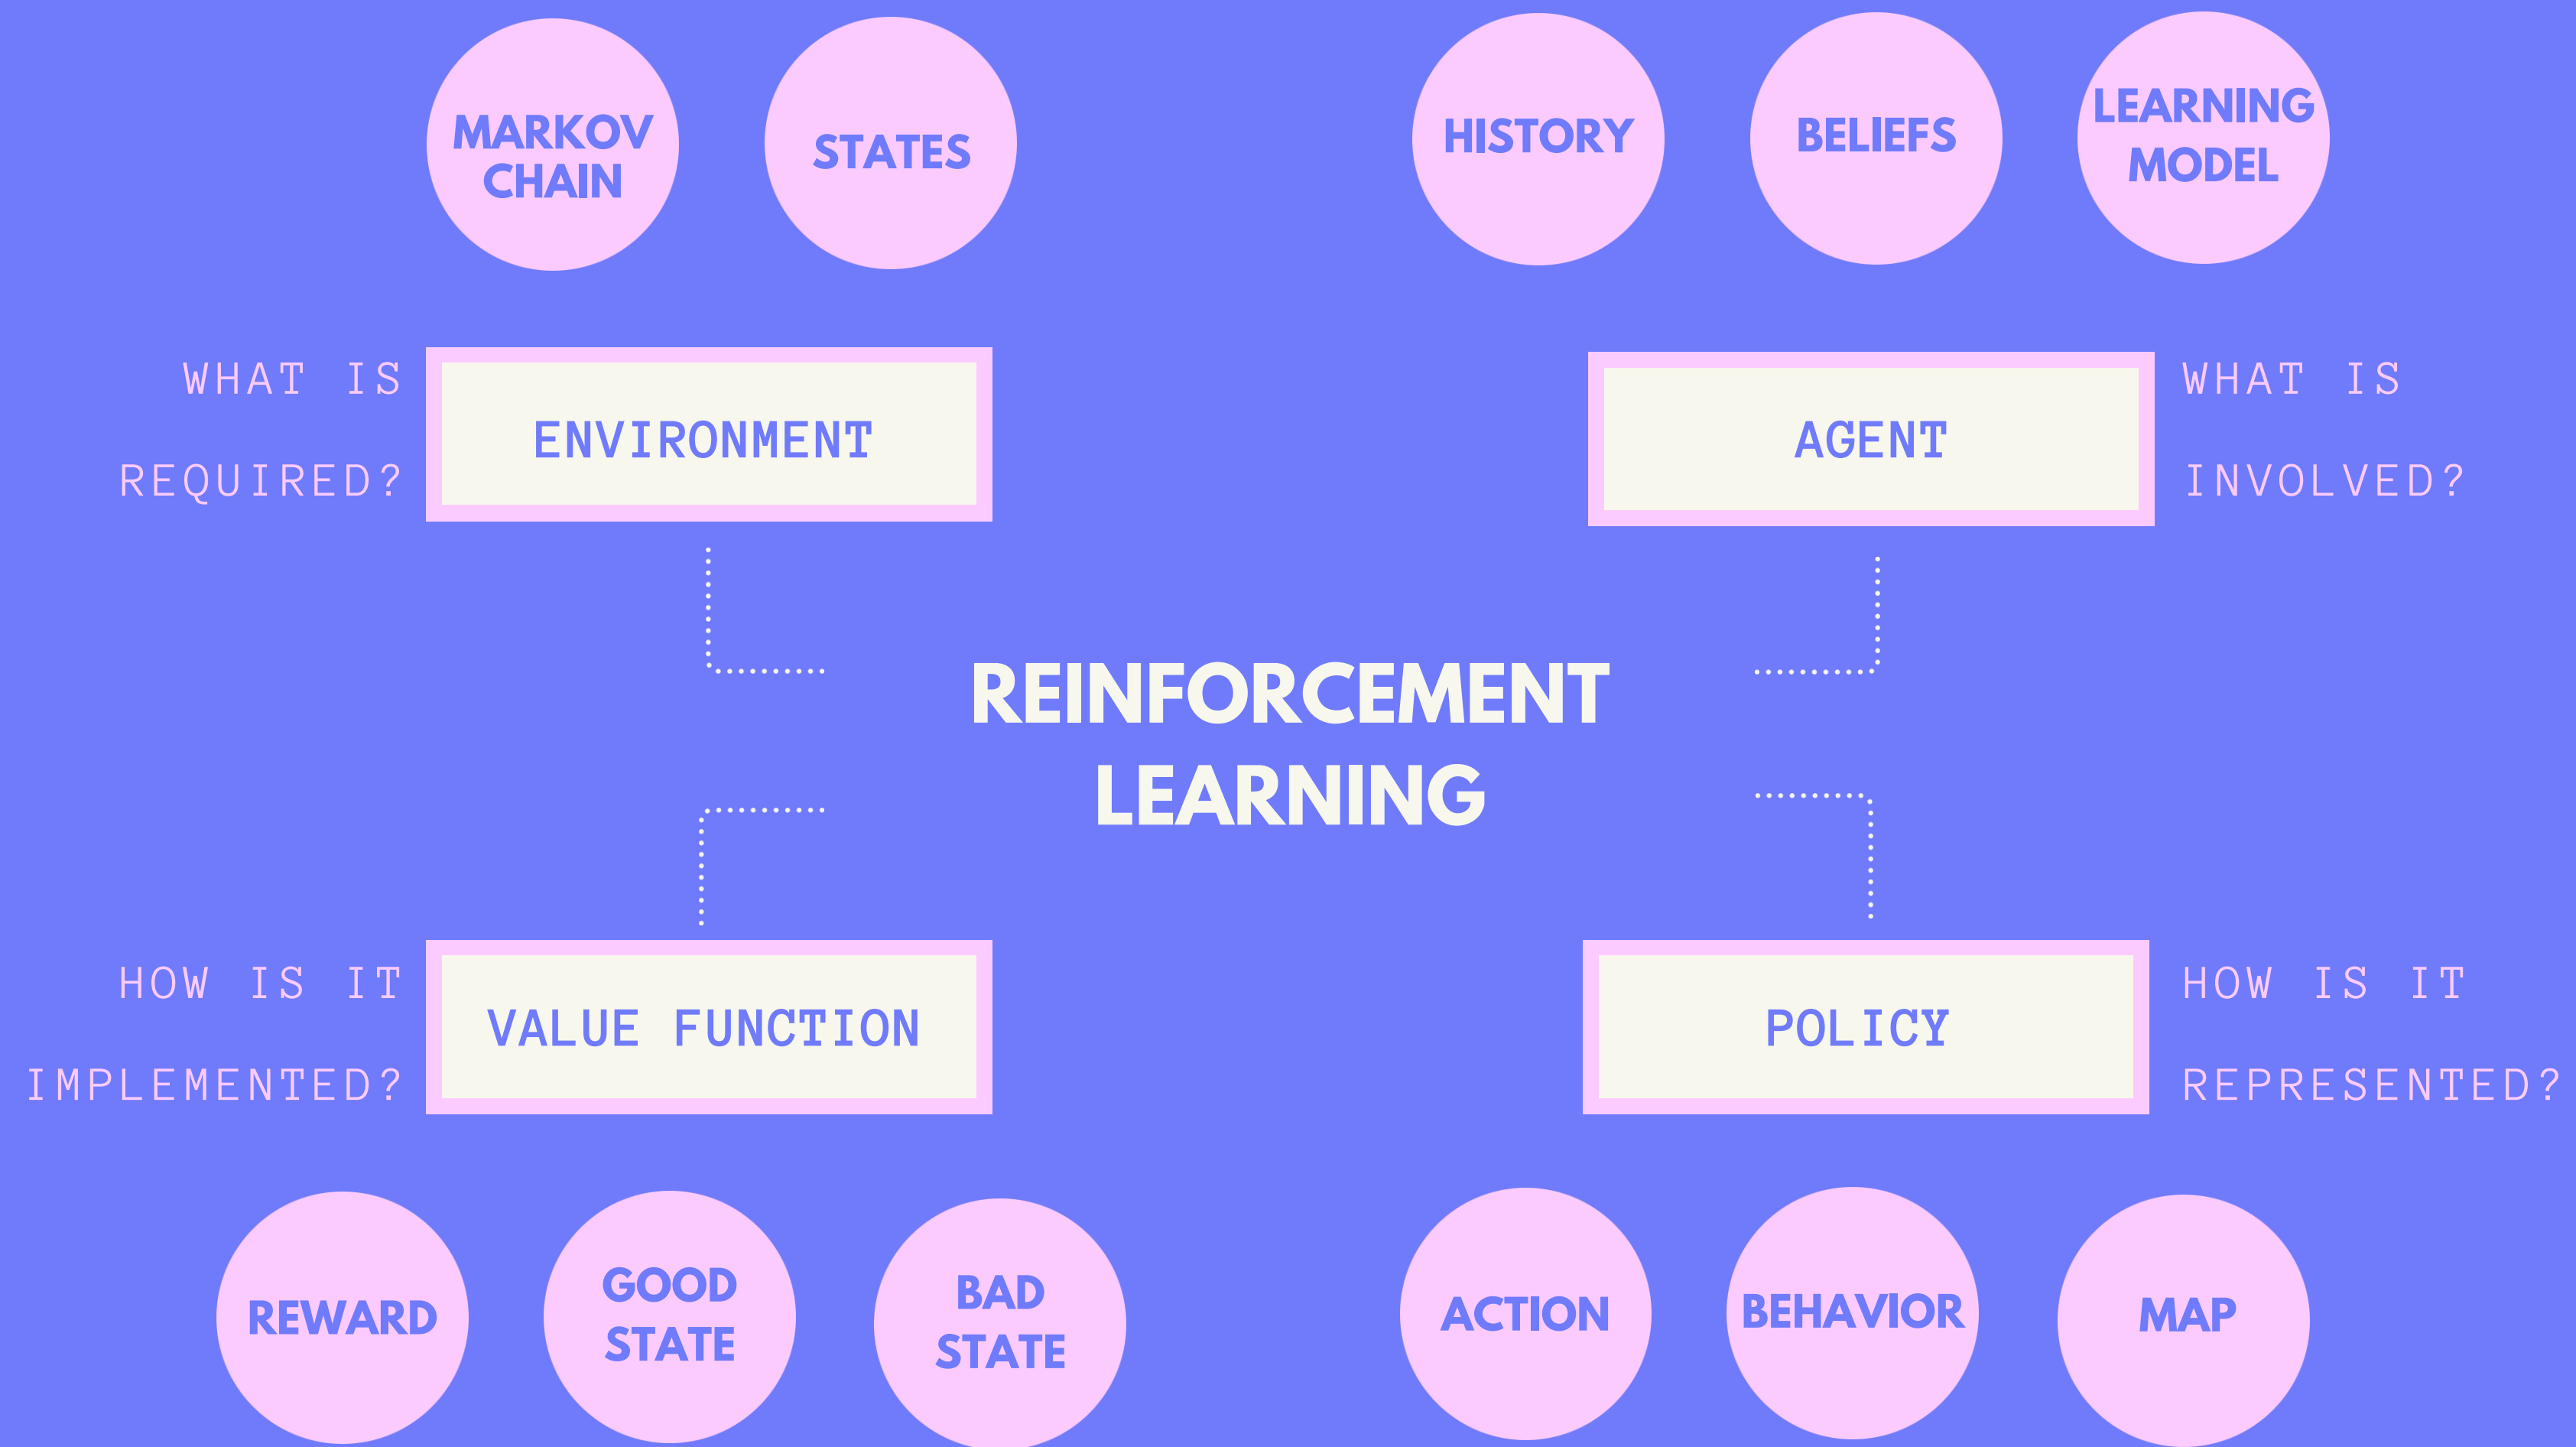
Kubernetes, Spark, Flask, Dask, etc

DATA MIRRORS HUMAN BIAS

<https://www.nytimes.com/2019/12/06/business/algorithm-bias-fix.html>

DATA AND HUMAN PRIVACY

<https://www.nytimes.com/series/new-york-times-privacy-project>




RL Use Cases

GAMING

AUTONOMOUS DRIVING

FINANCIAL TRADING

projects



shopping

```
def load_data(filename):
```

```
    """Follow the instructions in the docstring (the docstring  
    was too long to include on this slide)"""
```



shopping



```
def train_model(evidence, labels):  
    """  
    Given a list of evidence lists and a list of  
    labels, return a fitted k-nearest neighbor model  
    (k=1) trained on the data.  
    """
```

Check out the lecture if you are stuck!

shopping



```
def evaluate(labels, predictions):  
    """  
    Given a list of actual labels and a list of  
    predicted labels, return a tuple (sensitivity,  
    specificity). Assume each label is either a 1  
    (positive) or 0 (negative). `sensitivity` should be  
    a floating-point value from 0 to 1 representing the  
    "true positive rate": the proportion of actual  
    positive labels that were accurately identified.  
    `specificity` should be a floating-point value from  
    0 to 1 representing the "true negative rate": the  
    proportion of actual negative labels that were  
    accurately identified.  
    """
```

See next slide for hints!

confusion matrix

	Actual Positive	Actual Negative
Predicted Positive	True Positive	False Positive
Predicted Negative	False Negative	True Negative

sensitivity

	Actual Positive	Actual Negative
Predicted Positive	True Positive	False Positive
Predicted Negative	False Negative	True Negative

specificity

	Actual Positive	Actual Negative
Predicted Positive	True Positive	False Positive
Predicted Negative	False Negative	True Negative

sensitivity = (predicted and labeled positive)
/(labeled positive)


specificity = (predicted and labeled negative)
/(labeled negative)

sensitivity and **specificity** are two of many metrics to evaluate your machine learning model. always choose a metric that makes sense for your model and your business or research objective.


nim

```
def get_q_value(self, state, action):  
    """  
    Return the Q-value for the state `state` and the  
    action `action`. If no Q-value exists yet in  
    `self.q`, return 0.  
    """
```

Be mindful of the data type `state` and `action` should be




nim



```
def update_q_value(self, state, action, old_q, reward,
future_rewards):
    """
    Update the Q-value for the state `state` and the
    action `action` given the previous Q-value `old_q`,
    a current reward `reward`, and an estimate of
    future rewards `future_rewards`. Use the formula:
     $Q(s, a) \leftarrow \text{old value estimate} + \alpha * (\text{new value estimate} - \text{old value estimate})$ 
    where `old value estimate` is the previous Q-value, `alpha` is the
    learning rate, and `new value estimate` is the sum
    of the current reward and estimated future rewards.
    """
```

All that is expected is for you to implement the formula!


nim



```
def best_future_reward(self, state):  
    """  
    Given a state `state`, consider all possible  
    `(state, action)` pairs available in that state and  
    return the maximum of all their Q-values. Use 0 as  
    the Q-value if a `(state, action)` pair has no Q-  
    value in `self.q`. If there are no available  
    actions in `state`, return 0.  
    """
```

Hint: Q-values can be negative

nim



```
def choose_action(self, state, epsilon=True):  
    """  
    Given a state `state`, return an action `(i, j)` to  
    take. If `epsilon` is `False`, then return the best  
    action available in the state (the one with the  
    highest Q-value, using 0 for pairs that have no Q-  
    values). If `epsilon` is `True`, then with  
    probability `self.epsilon` choose a random  
    available action, otherwise choose the best action  
    available. If multiple actions have the same Q-  
    value, any of those options is an acceptable return  
    value.  
    """
```

The python module `random` will be helpful here! How do you determine if something is within a probability?