

Section with Barbara

Week 6

AI Stories



Projects




Language



A Friendly Reminder

**Project 6 is due this
Friday, August 7!**

AI Stories: IBM Watson



**Natural Language
Understanding
(NLU)**


AI can read

**Natural Language
Generation (NLG)**

AI can write

**Natural
Language
Processing (NLP)**

AI can analyze





*"Do you have anything between
10 am and 12 pm?"*



*"Depending on what service she would like.
What service is she looking for?"*

Current State of the Art in NLP: GPT-3

Artificial intelligence programs like deep learning neural networks may be able to beat humans at playing Go or chess, or doing arithmetic, or writing Navy Seal cypypasta, but they will never be able to truly think for themselves, to have consciousness, to feel any of the richness and complexity of the world that we mere humans can feel. Mere, unenlightened humans might be impressed by the abilities of simple deep learning programs, but when looked at in a more holistic manner, it all adds up to... well, nothing. They still don't exhibit any trace of consciousness. All of the available data support the notion that humans feel and experience the world differently than computers do. While a computer can beat a human master at chess or Go or some other game of structured rules, it will never be able to truly think outside of those rules, it will never be able to come up with its own new strategies on the fly, it will never be able to feel, to react, the way a human can. Artificial intelligence programs lack consciousness and self-awareness. They will never be able to have a sense of humor. They will never be able to appreciate art, or beauty, or love. They will never feel lonely. They will never have empathy for other people, for animals, for the environment. They will never enjoy music or fall in love, or cry at the drop of a hat. Merely by existing, mere, unenlightened humans are intellectually superior to computers, no matter how good our computers get at winning games like Go or Jeopardy. We don't live by the rules of those games. Our minds are much, much bigger than that. – <https://www.gwern.net/GPT-3>

Current State of the Art in NLP: GPT-3

GPT-3 wrote the text in the previous slide!

Language and AI

1950's: Beginnings,
Machine Translation

Pre-1990's: Knowledge
Bases and Rules

'90's and '00's:
Statistical Learning

2010's: Deep Learning

Now: Large Pretrained
Models

Deep Learning

2003: “A Neural Probabilistic Language Model”

2008: “A Unified Architecture for Natural Language Processing: Deep Neural Networks with Multitask Learning”


2013: “Efficient Estimation of Word Representations in Vector Space”

2014: “Sequence to Sequence Learning with Neural Networks”


2017: “Attention Is All You Need”

Word Embeddings

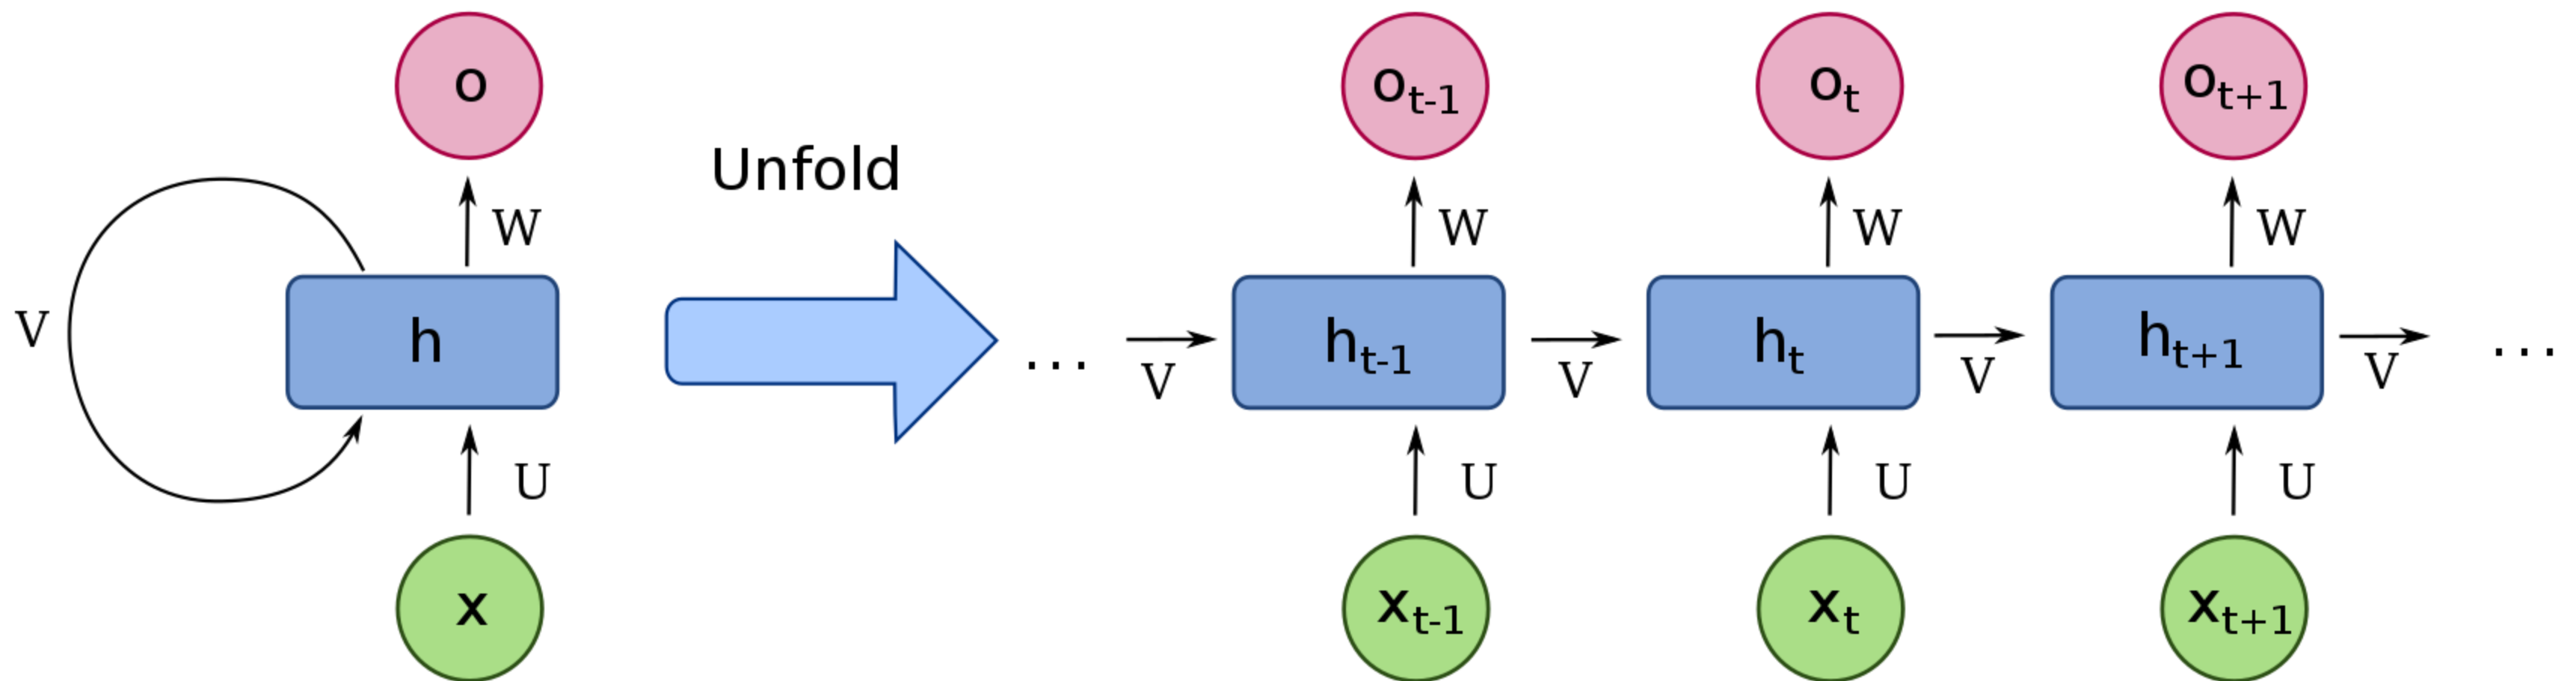
CBOW




Skip-Ngram

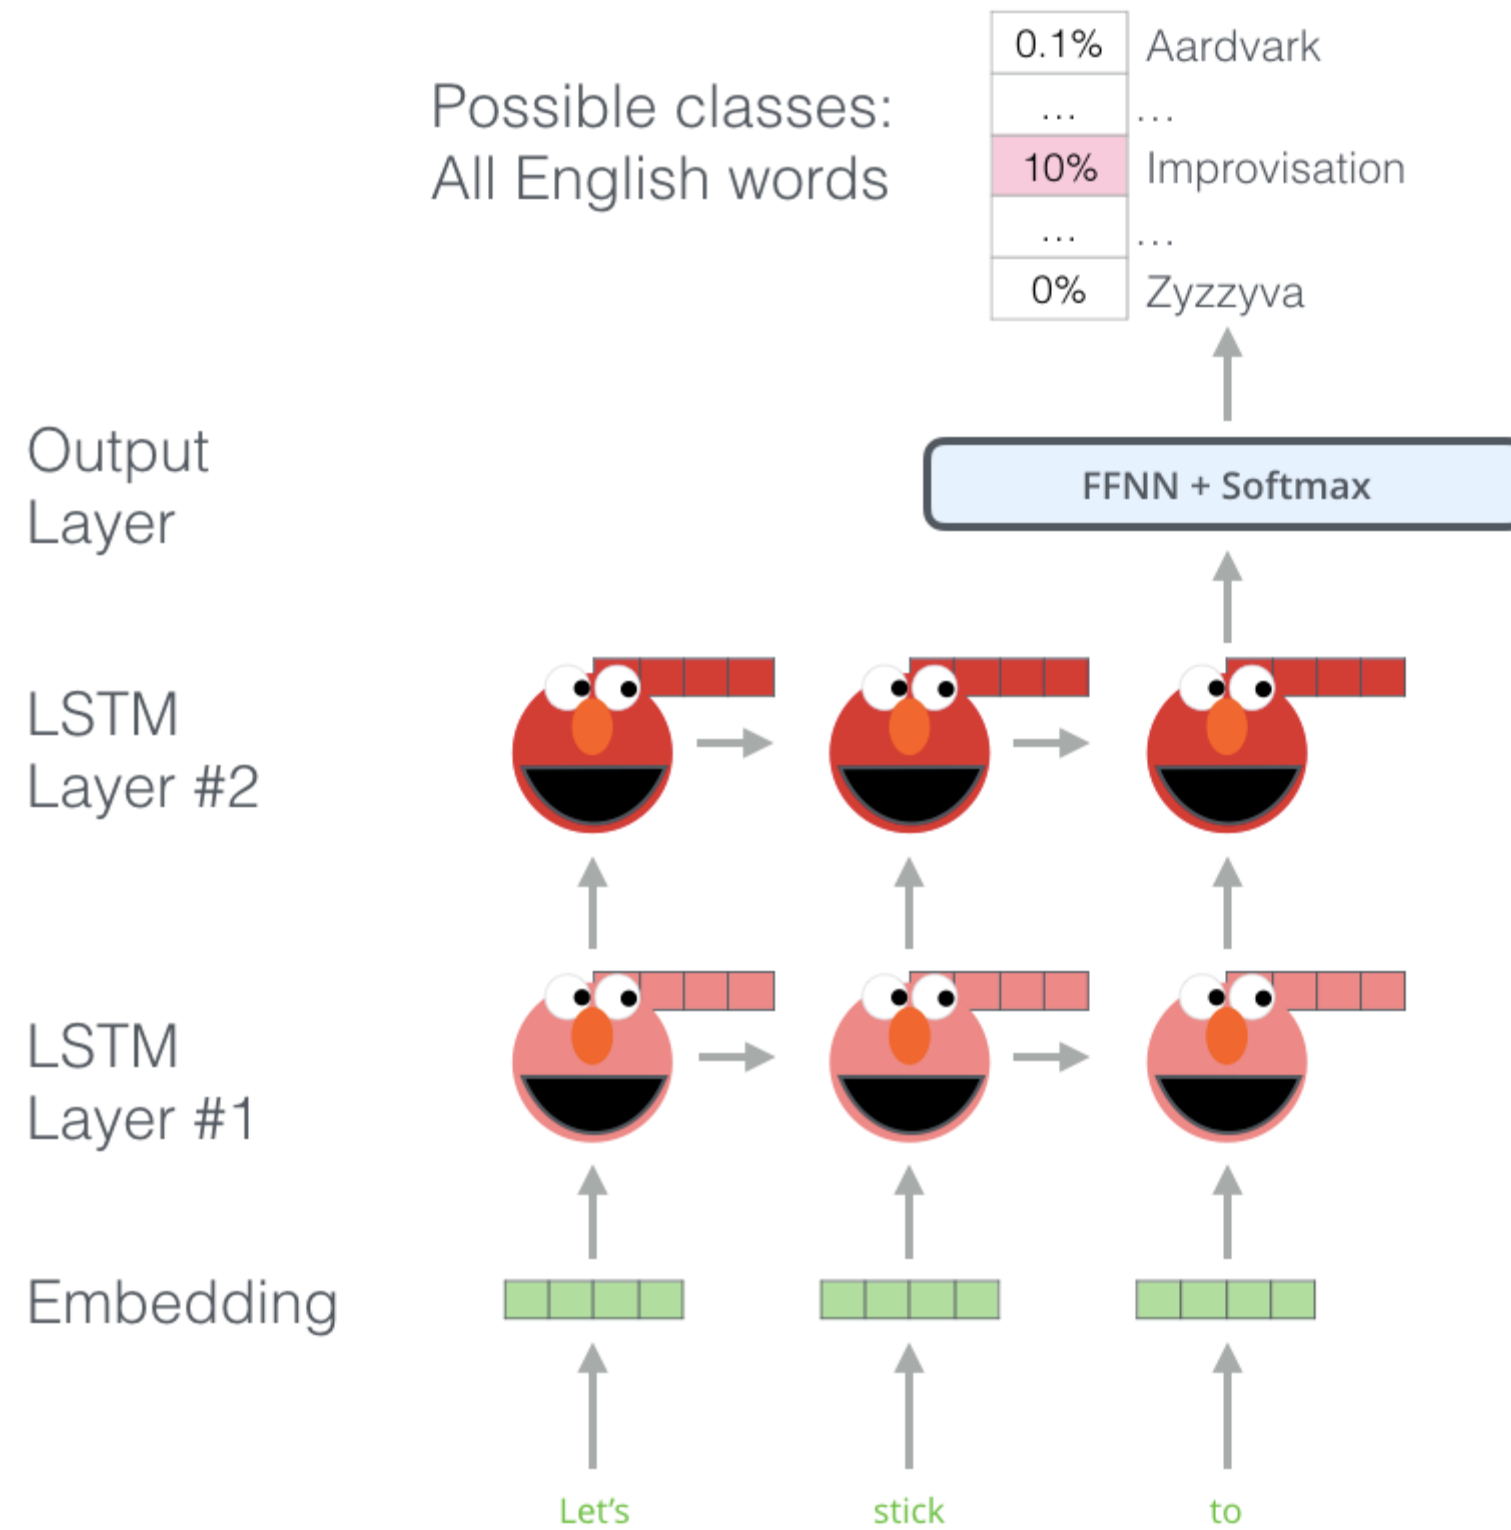


Recurrent Neural Network






ELMo – Bidirectional LSTM



seq2seq models


Encoder



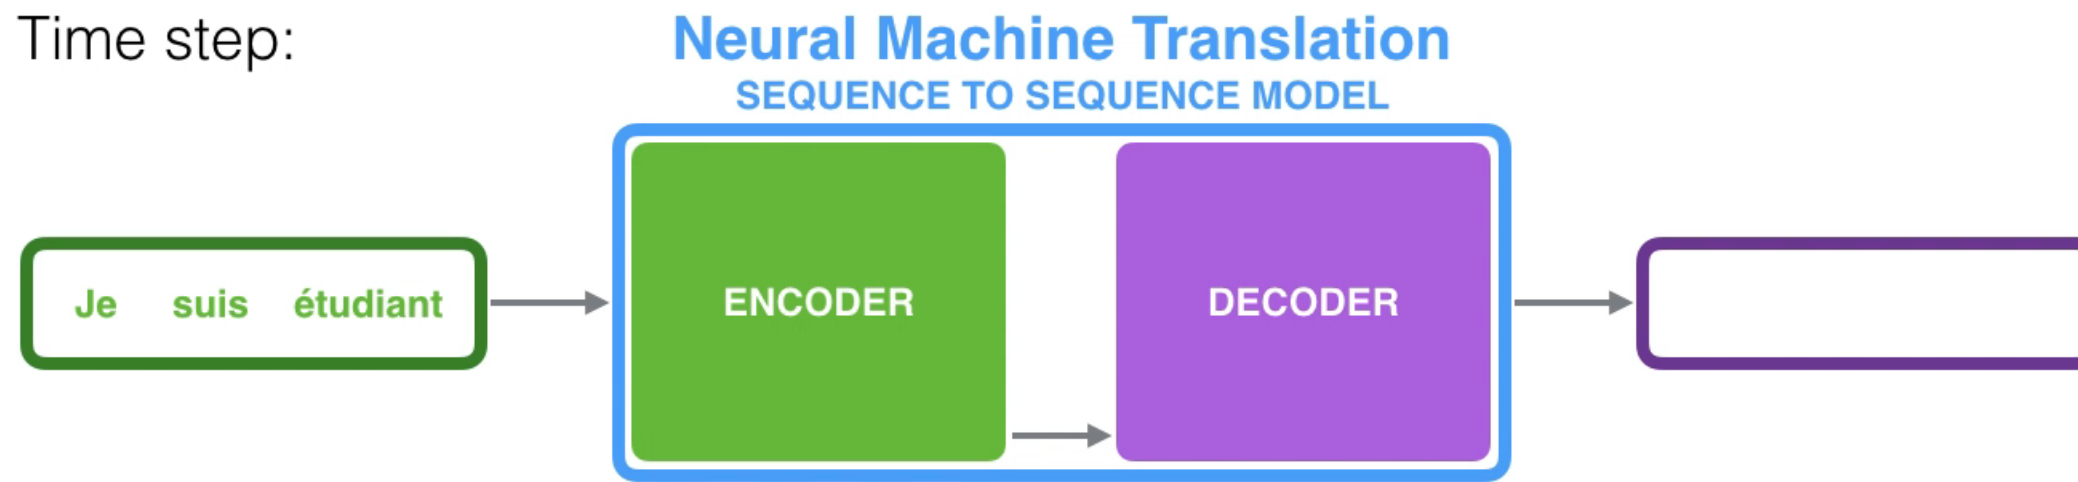
Context vector (length: 5)

$[0.1, -0.2, 0.8, 1.5, -0.3]$

Decoder

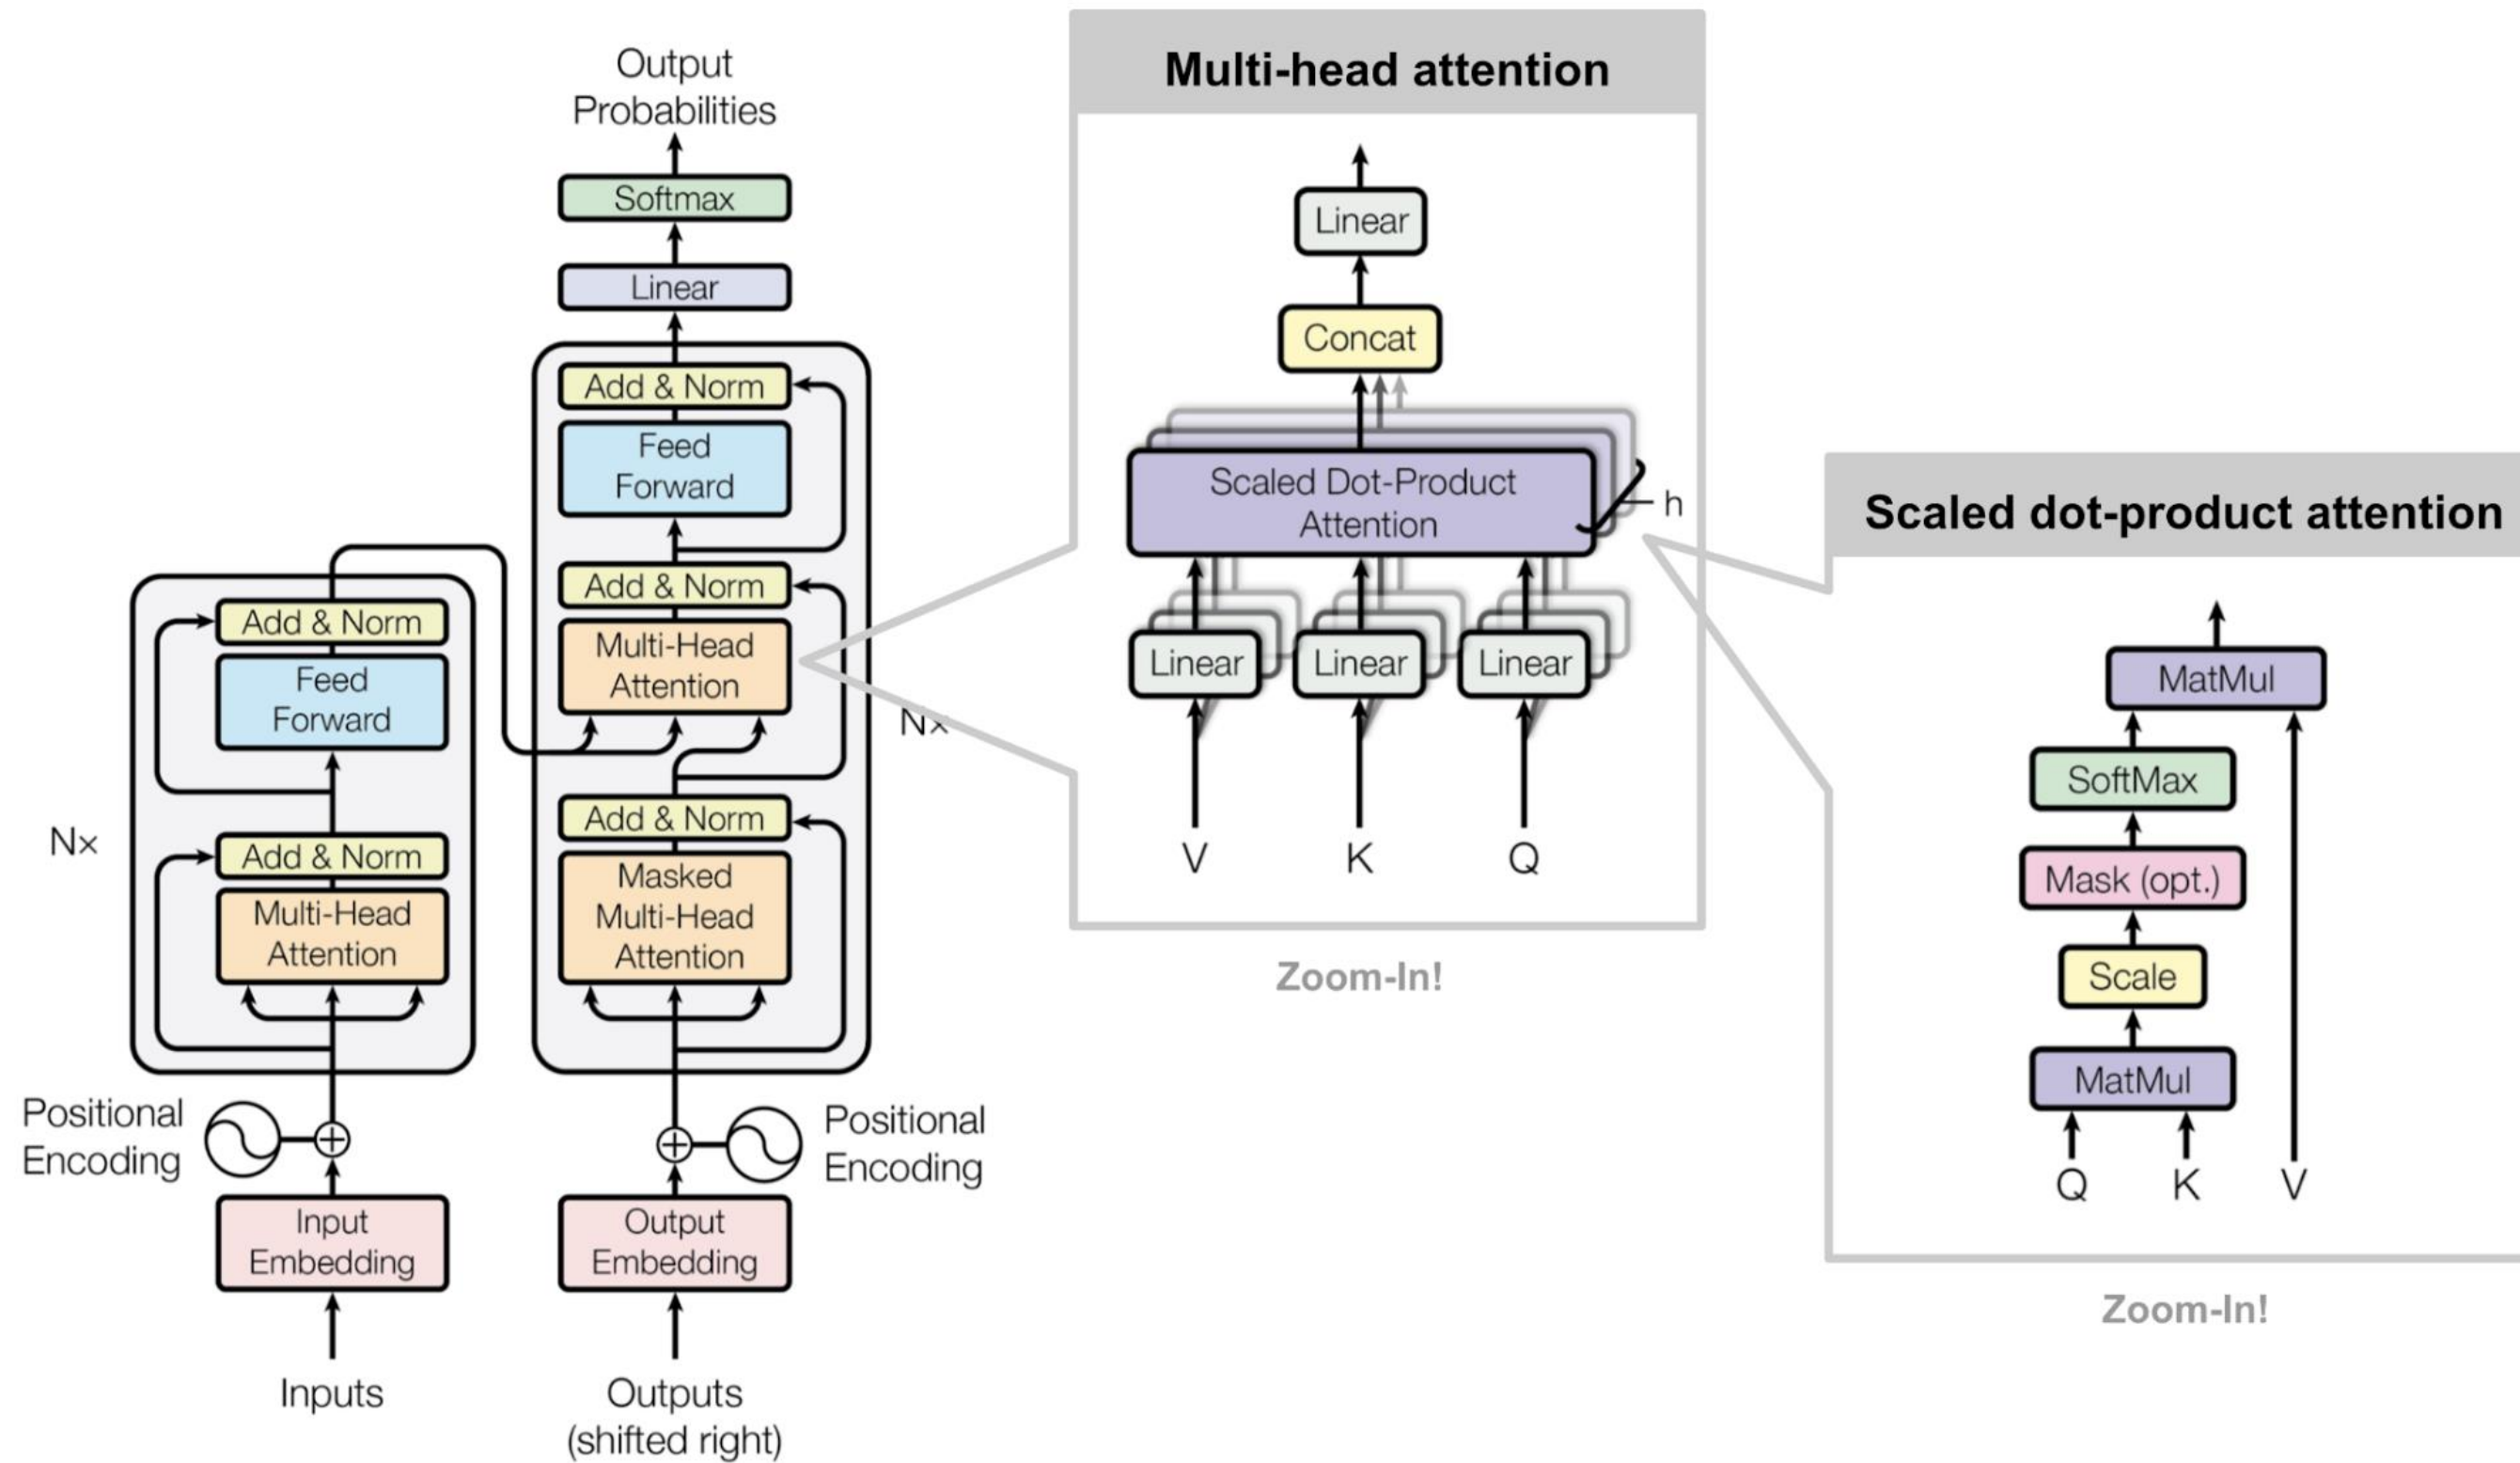


Time step:



Currently we are in the dawn of large, pretrained models able to perform extremely well with just a little bit of fine tuning.

Transformer





1) Concatenate all the attention heads



2) Multiply with a weight matrix W^o that was trained jointly with the model

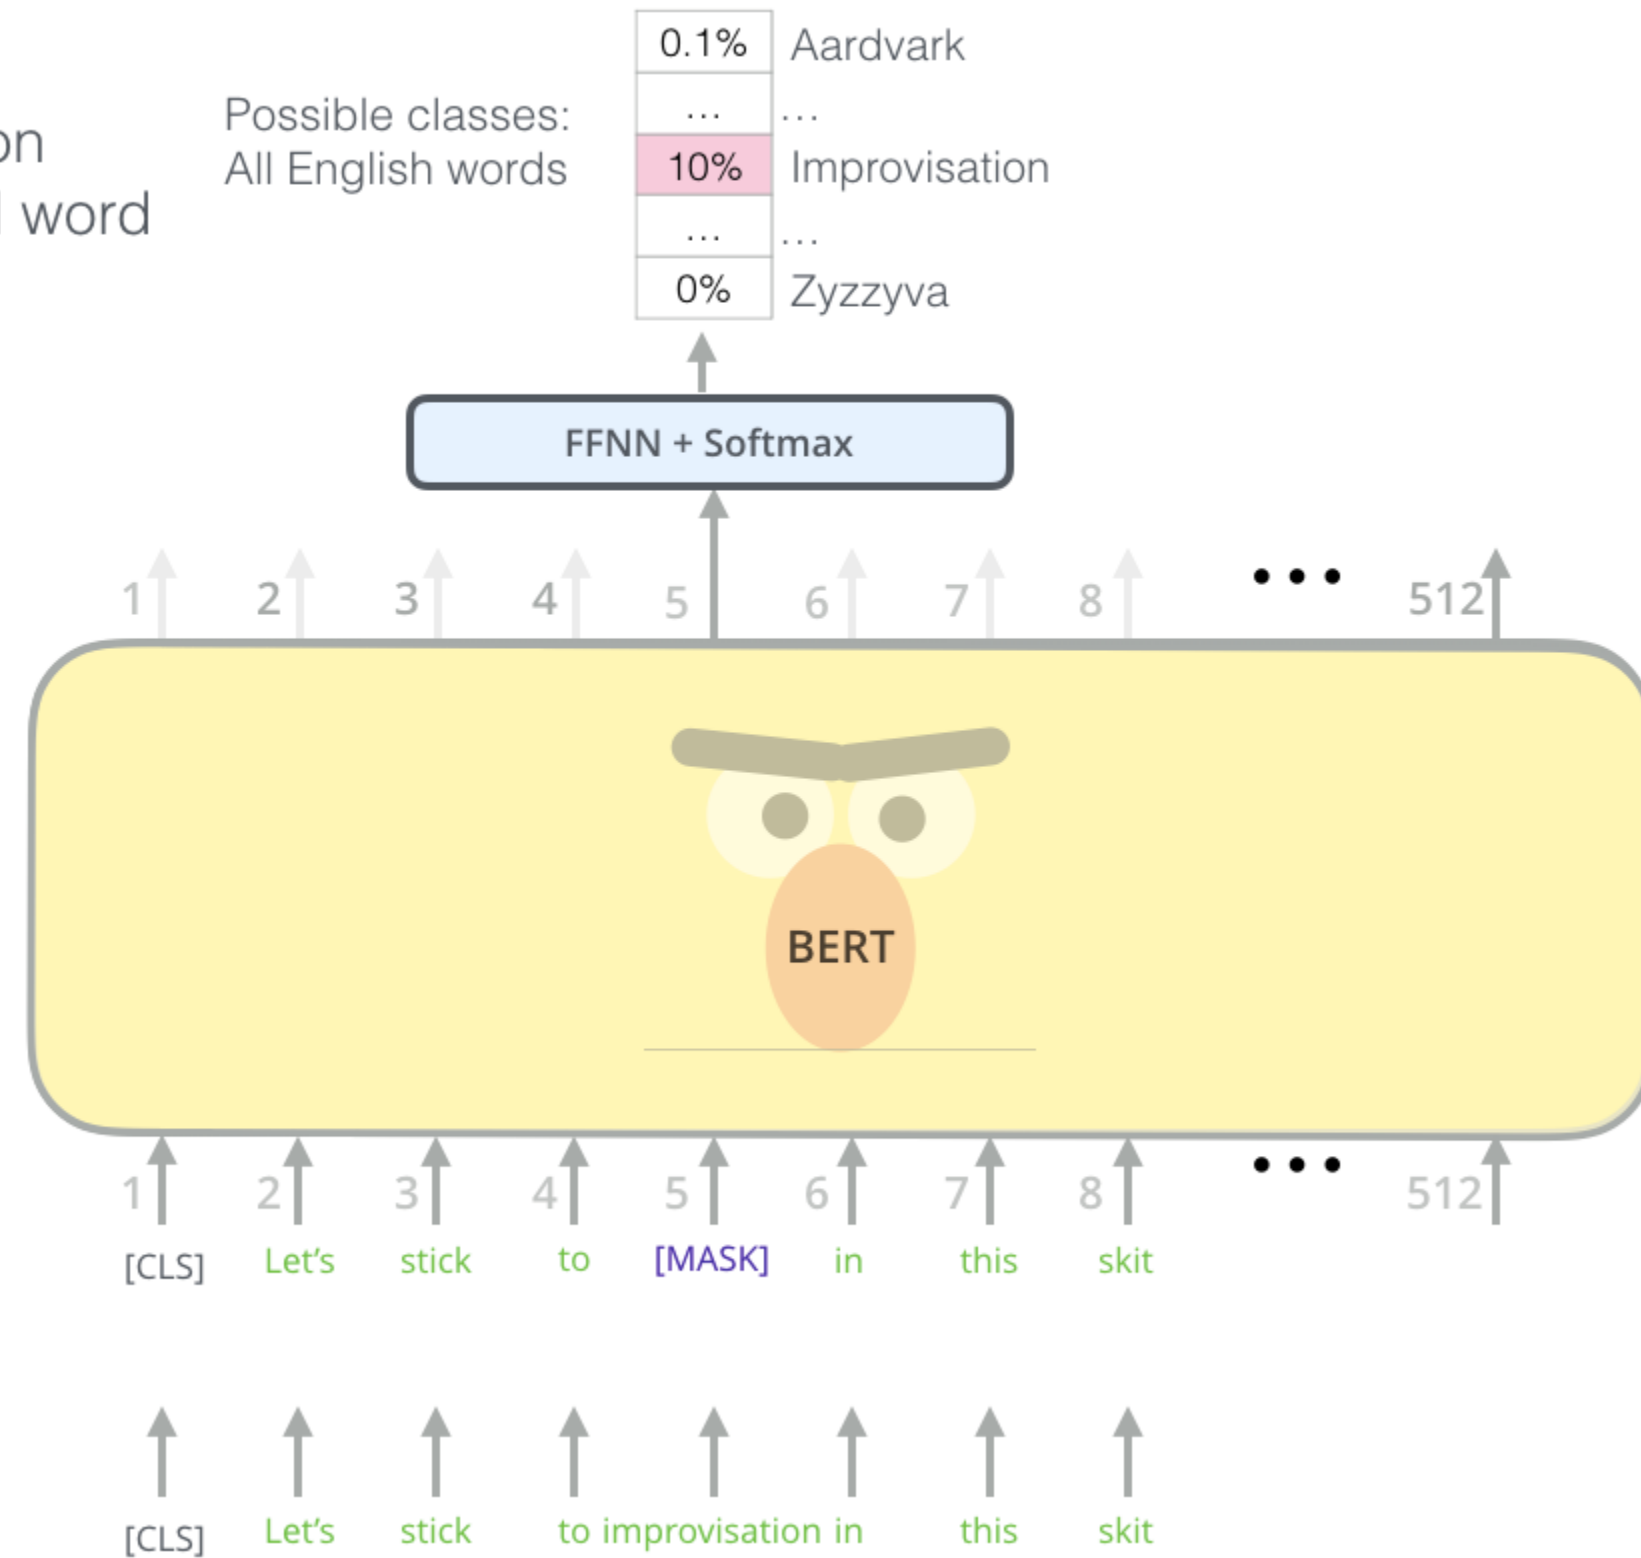
X

3) The result would be the **Z** matrix that captures information from all the attention heads. We can send this forward to the FFNN



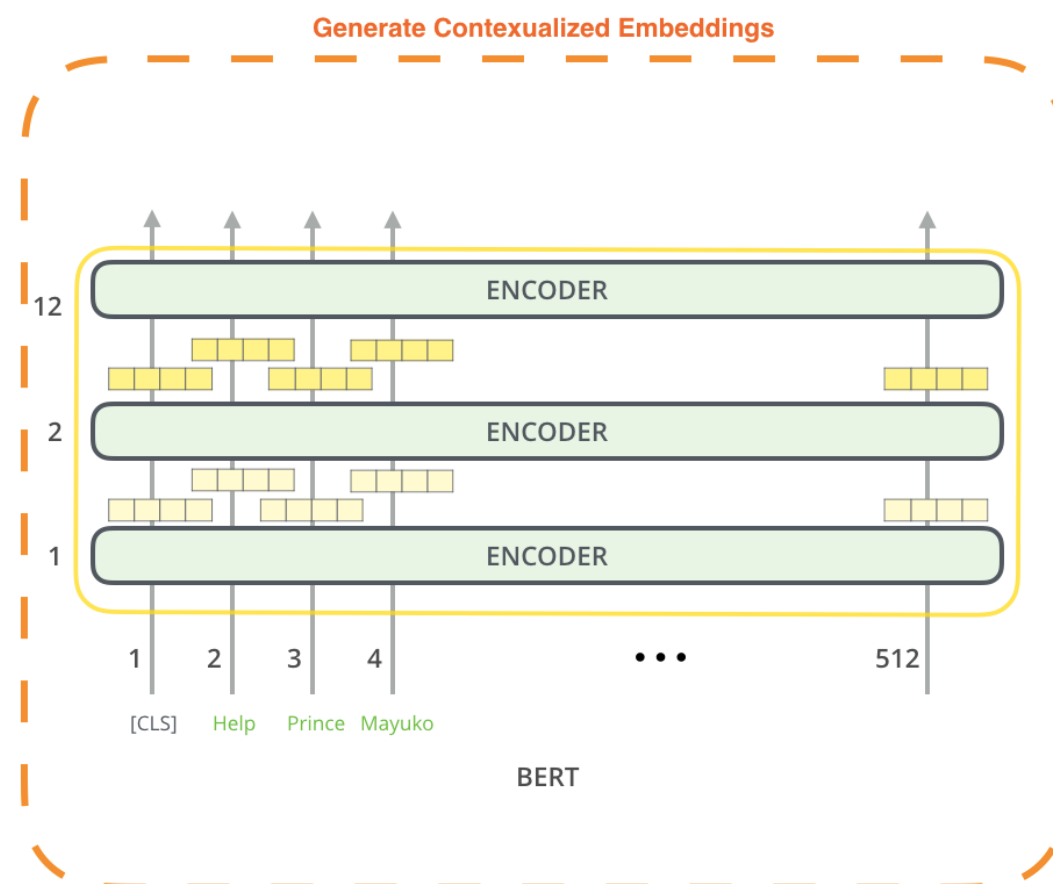
BERT

Use the output of the masked word's position to predict the masked word




Randomly mask 15% of tokens

Input



The output of each encoder layer along each token's path can be used as a feature representing that token.



But which one should we use?

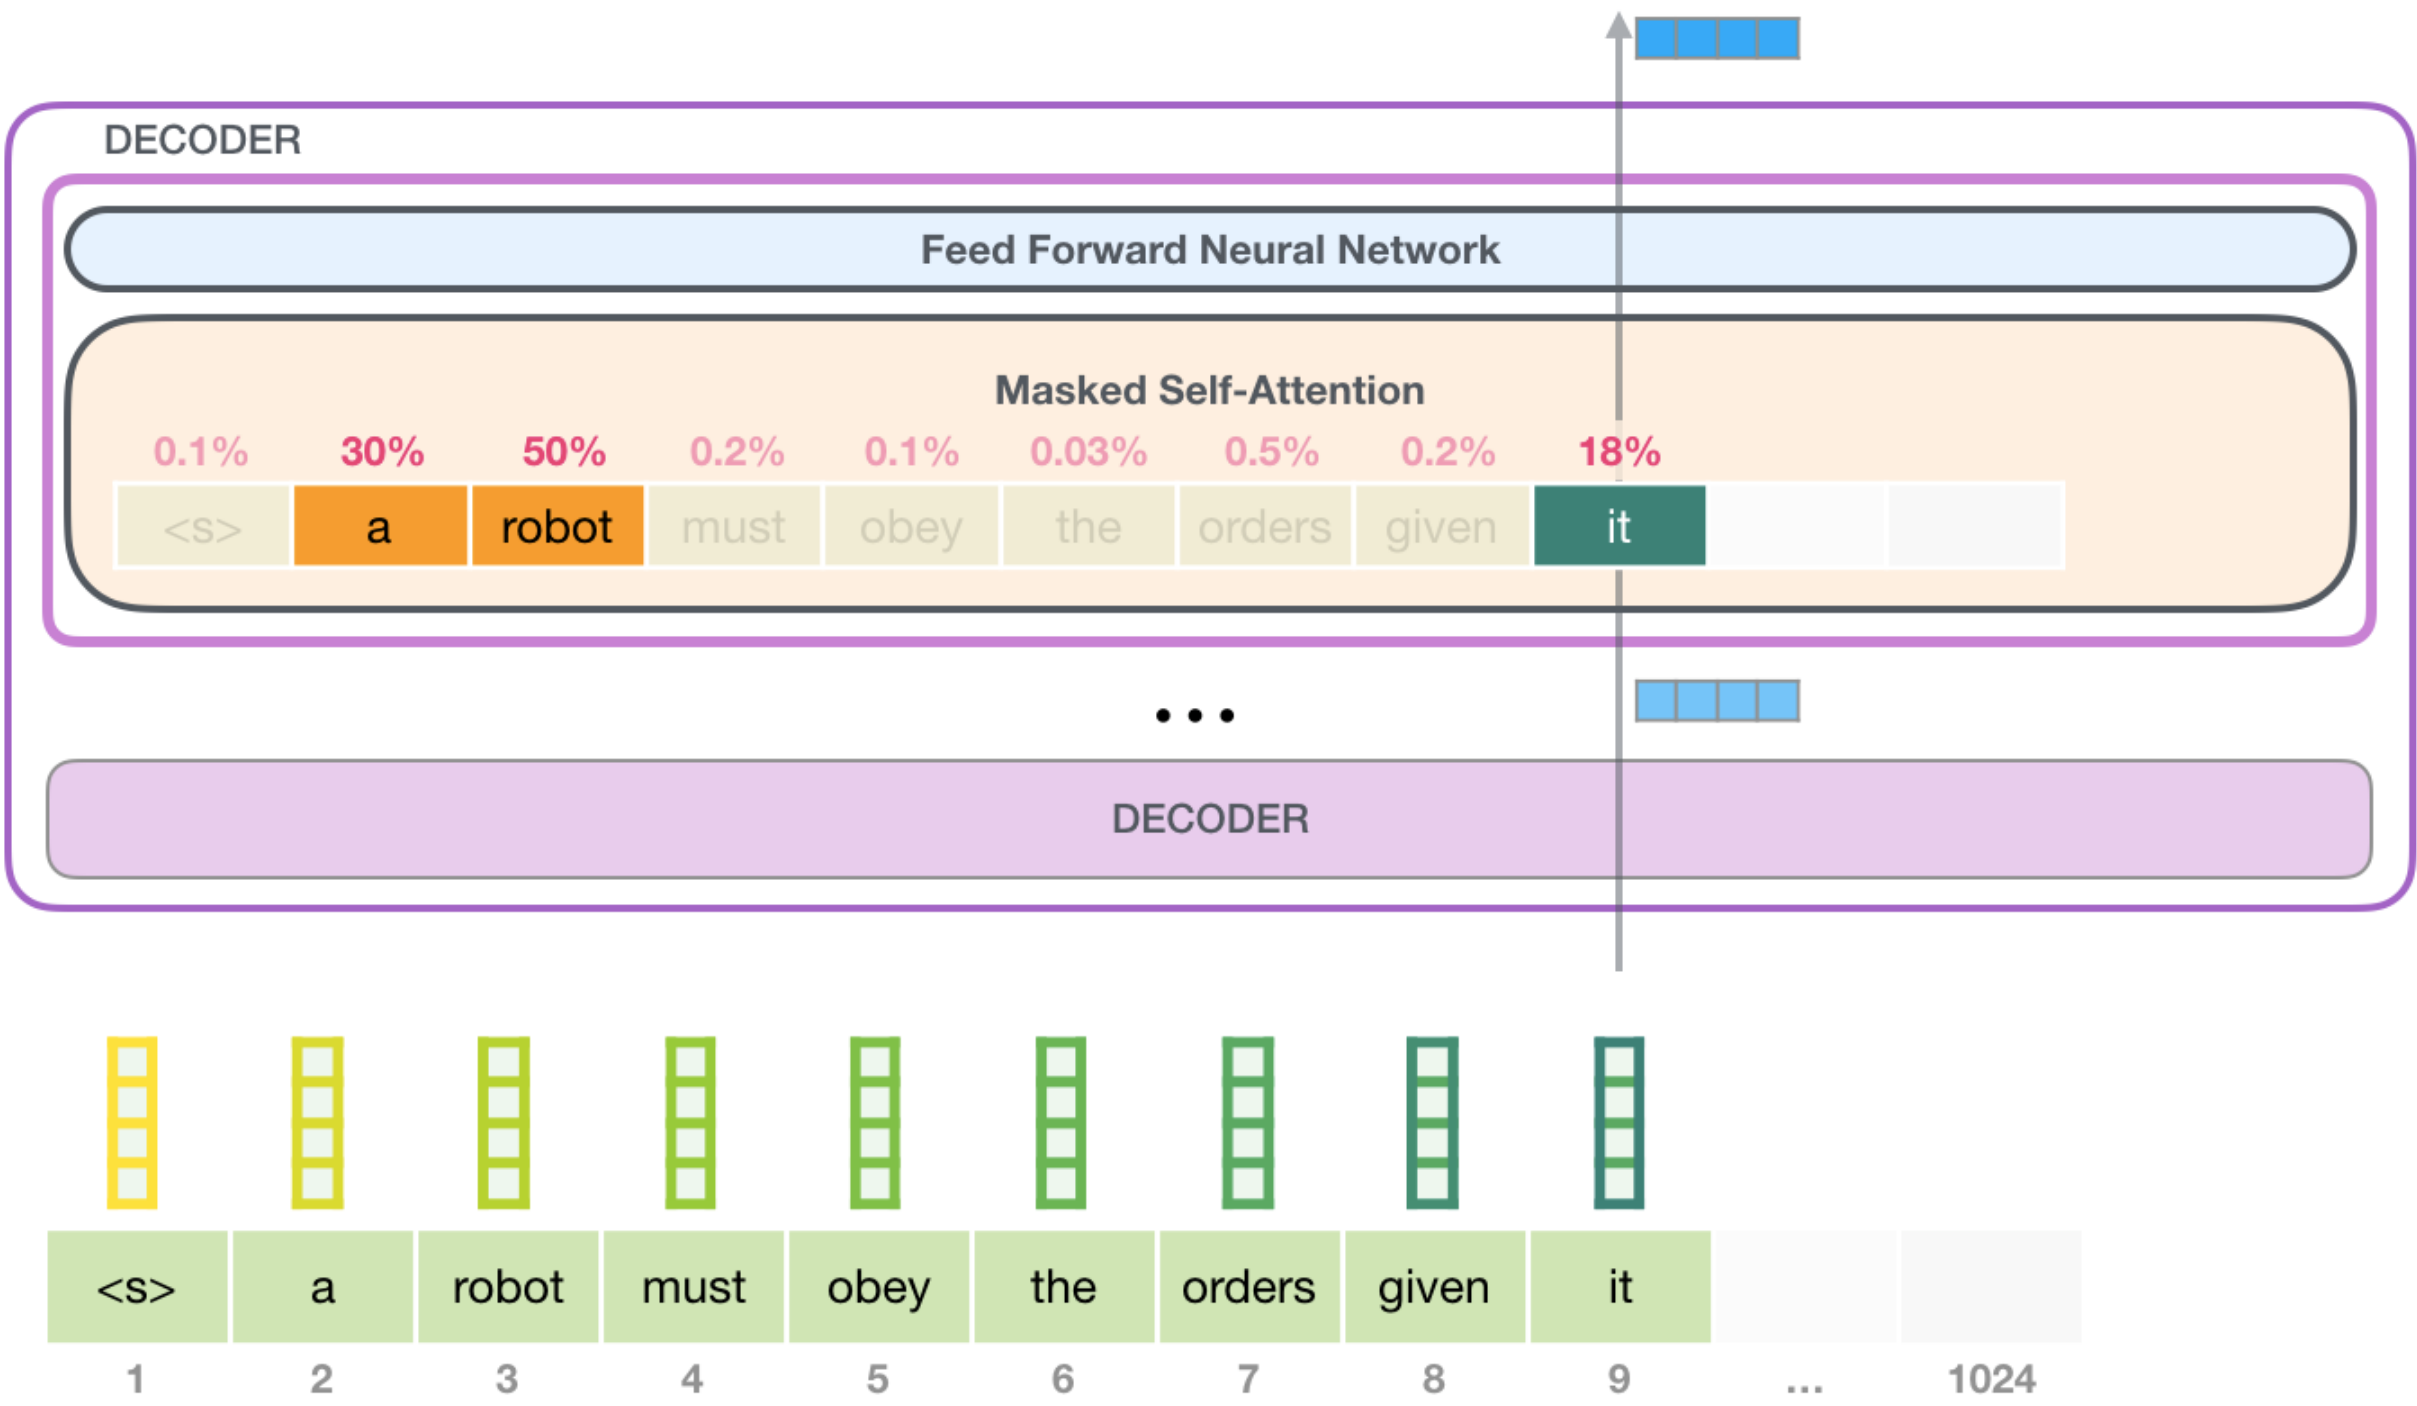
Context

We went to the *river* bank.

I need to go to bank to make a *deposit*.

Context

GPT



GPT-2

1.5 billion parameters

GPT-3

175 billion parameters.

Also got rid of gradient descent.

Reformer – the next state of the art?

[https://ai.googleblog.com/
2020/01/reformer-efficient-
transformer.html](https://ai.googleblog.com/2020/01/reformer-efficient-transformer.html)

projects



parser

```
NONTERMINALS = ""  
S -> N V  
""
```

The only grammar you need to model is the grammar in all the sentences.

The specifics of the design is up to you!



,

parser

```
def preprocess(sentence):  
    """  
    Convert `sentence` to a list of its words.  
    Pre-process sentence by converting all characters  
    to lowercase and removing any word that does not  
    contain at least one alphabetic character.  
    """
```

If you get stuck, refer to the lecture source code!



parser

```
def np_chunk(tree):  
    """  
    Return a list of all noun phrase chunks in the  
    sentence tree. A noun phrase chunk is defined as  
    any subtree of the sentence whose label is "NP"  
    that does not itself contain any other noun phrases  
    as subtrees.  
    """
```

`nlTK.tree` is all you need here! The `.subtree()` method
will be helpful here!



,

questions



```
def load_files(directory):  
    """
```

```
    Given a directory name, return a dictionary mapping  
    the filename of each `.txt` file inside that  
    directory to the file's contents as a string.  
    """
```

Remember to check for a file's `.txt` extension
before you import!

questions



```
def tokenize(document):  
    """  
    Given a document (represented as a string), return  
    a list of all the words in that document, in order.  
    Process document by converting all words to  
    lowercase and removing any punctuation or English  
    stopwords.  
    """
```

Hint: what case are the stopwords?

questions



```
def compute_idfs(documents):  
    """  
    Given a dictionary of `documents` that maps names  
    of documents to a list of words, return a  
    dictionary that maps words to their IDF values. Any  
    word that appears in at least one of the documents  
    should be in the resulting dictionary.  
    """
```

If you are stuck, refer to lecture source code!

questions



```
def top_files(query, files, idfs, n):
```

```
def top_sentences(query, sentences, idfs, n):
```

These are similar in implementation, but what is returned should be different! Hint, you can sum a list! The `str.count()` method can also be helpful here! Also, be mindful when you want words to repeat in your calculations and when you don't.

Yay! You Made It!

